

## 3 SQL and Pyspark Question 1: Identify VIP Users for Netflix

Question: To better cater to its most dedicated users, Netflix would like to identify its "VIP users" - those who are most active in terms of the number of hours of content they watch. Write a SQL query that will retrieve the top 10 users with the most watched hours in the last month.

Tables:

\* users table: user\_id (integer), sign\_up\_date (date), subscription\_type (text)

\* watching\_activity table: activity\_id (integer), user\_id (integer), date\_time (timestamp), show\_id (integer), hours\_watched (float)

%sql

```
select u.user_id, u.subscription_type, sum(wa.hours_watched) as total_hours_watched from users u join
watching_activity wa on u.user_id = wa.user_id where wa.data_time >= add_months(trunc(sysdate,'Month'), -1)
and wa.data_time < Trunc(sysdate,'Month') group by u.user_id, u.subscription_type order by total_hours_watched
desc fetch first 10 rows only;
```

%python

```
from pyspark.sql import functions as F
from pyspark.sql.window import Window
```

```
current_date = F.current_date()
first_day_of_month = F.date_trunc("month",current_date)
last_month_start = F.add_months(first_day_of_month, -1)
```

```
vip_users = (watching_activity_df.filter(F.col("date_time") >= last_month_start).filter(F.col("date_time") <
first_day_of_month).groupBy("user_id").agg(F.sum("hours_watched").alias("total_hours_watched")).join(users_df,"user
_id").orderBy(F.desc("total_hours_watched")).limit(10))
```

## 2 SQL and Pyspark Question 2: Analyzing Ratings For Netflix Shows

Question: Given a table of user ratings for Netflix shows, calculate the average rating for each show within a given month. Assume that there is a column for user\_id, show\_id, rating (out of 5 stars), and date of review. Order the results by month and then by average rating (descending order).

Tables:

\* show\_reviews table: review\_id (integer), user\_id (integer), review\_date (timestamp), show\_id (integer), stars (integer)

%sql

```
select show_id, extract(month from review_date) as month, extract(year from review_date) as year, avg(stars) as
avg_rating, count(*) as review_count from show_reviews group by show_id, extract(month from review_date),
extract(year from review_date) order by year,month,avg_rating desc;
```

```
%python
```

```
from pyspark.sql import functions as F
```

```
rating_by_month = (show_reviews_df.withColumn("month",
F.month("review_date")).withColumn("year",F.year("review_date")).groupBy("show_id","month","year").agg(F.avg("st
ars").alias("avg_rating"), F.count("*").alias("review_count")).orderBy("year","month",F.desc("avg_rating")))
rating_by_month.show()
```

SQL Question 3: What does EXCEPT / MINUS SQL commands do?

Question: Explain the purpose and usage of the EXCEPT (or MINUS in some SQL dialects) SQL commands.

The EXCEPT (or MINUS in Oracle) operator is used to return all distinct rows from the first query that are not present in the results of the second query.

Key characteristics:

Both queries must have the same number of columns

Corresponding columns must have compatible data types

Duplicates are removed from the final result

Similar to set difference operation in mathematics

```
%sql
```

```
-- Users who watched show A but not show B
```

```
SELECT user_id FROM watching_activity WHERE show_id = 101
```

```
EXCEPT
```

```
SELECT user_id FROM watching_activity WHERE show_id = 102;
```

```
%python
```

```
# Users who watched show A but not show B
```

```
show_a_watchers = watching_activity_df.filter(F.col("show_id") == 101).select("user_id")
```

```
show_b_watchers = watching_activity_df.filter(F.col("show_id") == 102).select("user_id")
```

```
a_not_b_watchers = show_a_watchers.exceptAll(show_b_watchers)
```

```
a_not_b_watchers.show()
```

SQL Question 4: Filter Netflix Users Based on Viewing History and Subscription Status

Question: You are given a database of Netflix's user viewing history and their current subscription status. Write a SQL query to find all active customers who watched more than 10 episodes of a show called "Stranger Things" in the last 30 days.

Tables:

- \* users table: user\_id (integer), active (boolean)

- \* viewing\_history table: user\_id (integer), show\_id (integer), episode\_id (integer), watch\_date (date)

- \* shows table: show\_id (integer), show\_name (text)

%sql

```
select u.user_id from users u join viewing_history vh on u.user_id = vh.user_id join shows s on vh.show_id = s.show_id
where u.active = True and s.show_name = 'Stranger Things' and vh.watch_date >= Sysdate - 30 group by u.user_id
having count(Distinct vh.episode_id) > 10;
```

%python

```
from pyspark.sql import functions as F
```

```
active_stranger_things_viewers = (users_df.filter(F.col("active") == True).join(viewing_history_df,
"user_id").join(shows_df.filter(F.col("show_name") == "Stranger Things"),
"show_id").groupBy("user_id").agg(F.countDistinct("episode_id").alias("episodes_watched"))).filter(F.col("episodes_watched") > 10))
active_stranger_things_viewers.show()
```

SQL Question 5: What does it mean to denormalize a database?

Question: Explain the concept and implications of denormalizing a database.

Denormalization is the process of intentionally introducing redundancy into a database design to improve read performance by reducing the number of joins needed for queries.

Key aspects:

Purpose: Optimize read performance at the cost of write performance

Techniques: Adding redundant columns, combining tables, creating summary tables

When to use: For read-heavy applications, reporting systems, data warehouses

Trade-offs:

Pros: Faster reads, simpler queries

Cons: Data redundancy, potential inconsistency, more complex updates

Increased storage requirements

In PySpark, denormalization typically involves:  
Creating wide DataFrames with redundant data  
Using join() operations to combine tables  
Persisting denormalized DataFrames for better performance

```
# Denormalize by joining user data with viewing history
denormalized_df = (viewing_history_df
    .join(users_df, "user_id")
    .join(shows_df, "show_id")
    .cache() # Persist in memory for faster access)
```

### SQL Question 6: Filter and Match Customer's Viewing Records

Question: As a data analyst at Netflix, you are asked to analyze the customer's viewing records. You confirmed that Netflix is especially interested in customers who have been continuously watching a particular genre - 'Documentary' over the last month. The task is to find the name and email of those customers who have viewed more than five 'Documentary' movies within the last month. 'Documentary' could be a part of a broader genre category in the genre field (for example, 'Documentary, History'). Therefore, the matching pattern could occur anywhere within the string.

Tables:

- \* movies table: movie\_id (integer), title (text), genre (text), release\_year (integer)
- \* customer table: user\_id (integer), name (text), email (text), last\_movie\_watched (integer), date\_watched (date)

```
%sql
select c.user_id, c.name, c.email, count(*) as documentary_count from customer c join movies m on
c.last_movie_watched = m.movie_id where c.date_watched >= sysdate - 30 and m.genre like '%Documentary%'
group by c.user_id, c.name, c.email having count(*) > 5 order by documentary_count DESC;
```

```
%python
from pyspark.sql import functions as F
```

```
documentary_viewers = (customer_df
    .join(movies_df, customer_df.last_movie_watched == movies_df.movie_id)
    .filter(F.col("date_watched") >= F.date_sub(F.current_date(), 30))
    .filter(F.col("genre").contains("Documentary"))
    .groupBy("user_id", "name", "email")
    .agg(F.count("*").alias("documentry_count"))
    .filter(F.col("documentry_count") > 5))
```

```
.orderBy(F.desc("documentry_count"))
```

```
documentray_viewers.show()
```